

```
In [10]: import warnings
```

```
In [11]: warnings.filterwarnings("ignore")
```

```
In [12]: import pandas as pd
```

```
In [13]: data=pd.read_csv("car.csv")
```

```
In [14]: #Display top 5 rows
```

```
In [15]: data.head()
```

Out[15]:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

```
In [16]: data.tail()
```

```
Out[16]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
296	city	2016	9.50	11.6	33988	Diesel	Dealer	Manual	0
297	brio	2015	4.00	5.9	60000	Petrol	Dealer	Manual	0
298	city	2009	3.35	11.0	87934	Petrol	Dealer	Manual	0
299	city	2017	11.50	12.5	9000	Diesel	Dealer	Manual	0
300	brio	2016	5.30	5.9	5464	Petrol	Dealer	Manual	0

```
In [17]: data.describe()
```

```
Out[17]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
count	301.000000	301.000000	301.000000	301.000000	301.000000
mean	2013.627907	4.661296	7.628472	36947.205980	0.043189
std	2.891554	5.082812	8.644115	38886.883882	0.247915
min	2003.000000	0.100000	0.320000	500.000000	0.000000
25%	2012.000000	0.900000	1.200000	15000.000000	0.000000
50%	2014.000000	3.600000	6.400000	32000.000000	0.000000
75%	2016.000000	6.000000	9.900000	48767.000000	0.000000
max	2018.000000	35.000000	92.600000	500000.000000	3.000000

```
In [18]: data.shape
```

```
Out[18]: (301, 9)
```

```
In [19]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null    object
1   Year            301 non-null    int64
2   Selling_Price   301 non-null    float64
3   Present_Price   301 non-null    float64
4   Kms_Driven      301 non-null    int64
5   Fuel_Type       301 non-null    object
6   Seller_Type     301 non-null    object
7   Transmission    301 non-null    object
8   Owner           301 non-null    int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [23]: data.isnull()
```

```
Out[23]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...
296	False	False	False	False	False	False	False	False	False
297	False	False	False	False	False	False	False	False	False
298	False	False	False	False	False	False	False	False	False
299	False	False	False	False	False	False	False	False	False
300	False	False	False	False	False	False	False	False	False

301 rows × 9 columns

```
In [24]: data.isnull().sum()
```

```
Out[24]: Car_Name      0
Year      0
Selling_Price  0
Present_Price  0
Kms_Driven   0
Fuel_Type    0
Seller_Type   0
Transmission  0
Owner        0
dtype: int64
```

```
In [21]: data.describe()
```

```
Out[21]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
count	301.000000	301.000000	301.000000	301.000000	301.000000
mean	2013.627907	4.661296	7.628472	36947.205980	0.043189
std	2.891554	5.082812	8.644115	38886.883882	0.247915
min	2003.000000	0.100000	0.320000	500.000000	0.000000
25%	2012.000000	0.900000	1.200000	15000.000000	0.000000
50%	2014.000000	3.600000	6.400000	32000.000000	0.000000
75%	2016.000000	6.000000	9.900000	48767.000000	0.000000
max	2018.000000	35.000000	92.600000	500000.000000	3.000000

```
In [22]: data.head(1)
```

```
Out[22]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0

```
In [25]: import datetime
```

```
In [27]: date_time=datetime.datetime.now()
```

```
In [30]: data["Age"]=date_time.year-data["Year"]
```

```
In [31]: data.head()
```

```
Out[31]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0	10
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0	11
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0	7
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0	13
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0	10

```
In [32]: data.drop("Year",axis=1,inplace=True)
```

```
In [33]: data
```


```
Out[33]:
```

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	Petrol	Dealer	Manual	0	10
1	sx4	4.75	9.54	43000	Diesel	Dealer	Manual	0	11
2	ciaz	7.25	9.85	6900	Petrol	Dealer	Manual	0	7
3	wagon r	2.85	4.15	5200	Petrol	Dealer	Manual	0	13
4	swift	4.60	6.87	42450	Diesel	Dealer	Manual	0	10
...
296	city	9.50	11.60	33988	Diesel	Dealer	Manual	0	8
297	brio	4.00	5.90	60000	Petrol	Dealer	Manual	0	9
298	city	3.35	11.00	87934	Petrol	Dealer	Manual	0	15
299	city	11.50	12.50	9000	Diesel	Dealer	Manual	0	7
300	brio	5.30	5.90	5464	Petrol	Dealer	Manual	0	8

301 rows × 9 columns

```
In [34]: import seaborn as sns
```

```
In [35]: sns.boxplot(data["Selling_Price"])
```



```
In [39]: data[["Selling_Price"]]
```


```
Out[39]:
```

	Selling_Price
0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
...	...
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

301 rows × 1 columns



```
In [36]: sns.boxplot(data["Kms_Driven"])
```

```
Out[36]: <Axes: >
```



```
In [37]: sns.boxplot(data["Age"])
```

```
Out[37]: <Axes: >
```



```
In [40]: sorted(data["Selling_Price"],reverse=True)
```

```
Out[40]: [35.0,  
          33.0,  
          23.5,  
          23.0,  
          23.0,  
          23.0,  
          20.75,  
          19.99,  
          19.75,  
          18.75,  
          18.0,  
          17.0,  
          16.0,  
          14.9,  
          14.73,  
          14.5,  
          14.25,  
          12.9,  
          12.5,  
          11.75]
```

```
In [43]: data=data[~(data["Selling_Price"]>=33.0) & (data["Selling_Price"]<=35.0)]
```

```
In [44]: data
```

```
Out[44]:
```

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	Petrol	Dealer	Manual	0	10
1	sx4	4.75	9.54	43000	Diesel	Dealer	Manual	0	11
2	ciaz	7.25	9.85	6900	Petrol	Dealer	Manual	0	7
3	wagon r	2.85	4.15	5200	Petrol	Dealer	Manual	0	13
4	swift	4.60	6.87	42450	Diesel	Dealer	Manual	0	10
...
296	city	9.50	11.60	33988	Diesel	Dealer	Manual	0	8
297	brio	4.00	5.90	60000	Petrol	Dealer	Manual	0	9
298	city	3.35	11.00	87934	Petrol	Dealer	Manual	0	15
299	city	11.50	12.50	9000	Diesel	Dealer	Manual	0	7
300	brio	5.30	5.90	5464	Petrol	Dealer	Manual	0	8

299 rows × 9 columns

```
In [45]: #Encoding the categorical Columns
```

```
In [47]: data.head(1)
```

```
Out[47]:
```

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	Petrol	Dealer	Manual	0	10

```
In [49]: data["Fuel_Type"].unique()
```

```
Out[49]: array(['Petrol', 'Diesel', 'CNG'], dtype=object)
```

```
In [50]: data["Fuel_Type"]=data["Fuel_Type"].map({"Petrol":0,"Diesel":1,"CNG":2})
```

```
In [51]: data
```

```
Out[51]:
```

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	0	Dealer	Manual	0	10
1	sx4	4.75	9.54	43000	1	Dealer	Manual	0	11
2	ciaz	7.25	9.85	6900	0	Dealer	Manual	0	7
3	wagon r	2.85	4.15	5200	0	Dealer	Manual	0	13
4	swift	4.60	6.87	42450	1	Dealer	Manual	0	10
...
296	city	9.50	11.60	33988	1	Dealer	Manual	0	8
297	brio	4.00	5.90	60000	0	Dealer	Manual	0	9
298	city	3.35	11.00	87934	0	Dealer	Manual	0	15
299	city	11.50	12.50	9000	1	Dealer	Manual	0	7
300	brio	5.30	5.90	5464	0	Dealer	Manual	0	8

299 rows × 9 columns

```
In [52]: data["Fuel_Type"].unique()
```

```
Out[52]: array([0, 1, 2], dtype=int64)
```

```
In [53]: data["Seller_Type"].unique()
```

```
Out[53]: array(['Dealer', 'Individual'], dtype=object)
```

```
In [54]: data["Seller_Type"]=data["Seller_Type"].map({"Dealer":0,"Individual":1})
```

```
In [55]: data
```

```
Out[55]:
```

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	0	0	Manual	0	10
1	sx4	4.75	9.54	43000	1	0	Manual	0	11
2	ciaz	7.25	9.85	6900	0	0	Manual	0	7
3	wagon r	2.85	4.15	5200	0	0	Manual	0	13
4	swift	4.60	6.87	42450	1	0	Manual	0	10
...
296	city	9.50	11.60	33988	1	0	Manual	0	8
297	brio	4.00	5.90	60000	0	0	Manual	0	9
298	city	3.35	11.00	87934	0	0	Manual	0	15
299	city	11.50	12.50	9000	1	0	Manual	0	7
300	brio	5.30	5.90	5464	0	0	Manual	0	8

299 rows × 9 columns

```
In [56]: data["Transmission"].unique()
```

```
Out[56]: array(['Manual', 'Automatic'], dtype=object)
```

```
In [57]: data["Transmission"]=data["Transmission"].map({"Manual":0,"Automatic":1})
```

In [58]: data

Out[58]:

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	0	0	0	0	10
1	sx4	4.75	9.54	43000	1	0	0	0	11
2	ciaz	7.25	9.85	6900	0	0	0	0	7
3	wagon r	2.85	4.15	5200	0	0	0	0	13
4	swift	4.60	6.87	42450	1	0	0	0	10
...
296	city	9.50	11.60	33988	1	0	0	0	8
297	brio	4.00	5.90	60000	0	0	0	0	9
298	city	3.35	11.00	87934	0	0	0	0	15
299	city	11.50	12.50	9000	1	0	0	0	7
300	brio	5.30	5.90	5464	0	0	0	0	8

299 rows × 9 columns


```
In [59]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 299 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        299 non-null    object
1   Selling_Price    299 non-null    float64
2   Present_Price    299 non-null    float64
3   Kms_Driven       299 non-null    int64
4   Fuel_Type        299 non-null    int64
5   Seller_Type      299 non-null    int64
6   Transmission     299 non-null    int64
7   Owner            299 non-null    int64
8   Age              299 non-null    int64
dtypes: float64(2), int64(6), object(1)
memory usage: 23.4+ KB
```



```
In [60]: sns.boxplot(data["Kms_Driven"])
```

```
Out[60]: <Axes: >
```



```
In [67]: data=data[~(data["Kms_Driven"]>100000)]
```

In [68]: data


Out[68]:

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	0	0	0	0	10
1	sx4	4.75	9.54	43000	1	0	0	0	11
2	ciaz	7.25	9.85	6900	0	0	0	0	7
3	wagon r	2.85	4.15	5200	0	0	0	0	13
4	swift	4.60	6.87	42450	1	0	0	0	10
...
296	city	9.50	11.60	33988	1	0	0	0	8
297	brio	4.00	5.90	60000	0	0	0	0	9
298	city	3.35	11.00	87934	0	0	0	0	15
299	city	11.50	12.50	9000	1	0	0	0	7
300	brio	5.30	5.90	5464	0	0	0	0	8

291 rows × 9 columns

```
In [69]: sns.boxplot(data["Kms_Driven"])
```

```
Out[69]: <Axes: >
```



```
In [70]: X=data.drop(["Car_Name","Selling_Price"],axis=1)
```

```
In [71]: y=data["Selling_Price"]
```

In [72]: X

Out[72]:

	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	5.59	27000	0	0	0	0	10
1	9.54	43000	1	0	0	0	11
2	9.85	6900	0	0	0	0	7
3	4.15	5200	0	0	0	0	13
4	6.87	42450	1	0	0	0	10
...
296	11.60	33988	1	0	0	0	8
297	5.90	60000	0	0	0	0	9
298	11.00	87934	0	0	0	0	15
299	12.50	9000	1	0	0	0	7
300	5.90	5464	0	0	0	0	8

291 rows × 7 columns

In [74]: y

Out[74]:

```
0      3.35
1      4.75
2      7.25
3      2.85
4      4.60
...
296    9.50
297    4.00
298    3.35
299   11.50
300    5.30
```

Name: Selling_Price, Length: 291, dtype: float64

```
In [75]: from sklearn.model_selection import train_test_split
```

```
In [76]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.20,random_state=42)
```

```
In [77]: data.head()
```

Out[77]:

	Car_Name	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Age
0	ritz	3.35	5.59	27000	0	0	0	0	10
1	sx4	4.75	9.54	43000	1	0	0	0	11
2	ciaz	7.25	9.85	6900	0	0	0	0	7
3	wagon r	2.85	4.15	5200	0	0	0	0	13
4	swift	4.60	6.87	42450	1	0	0	0	10

```
In [80]: from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from xgboost import XGBRegressor
```

```
In [81]: lr=LinearRegression()
lr.fit(X_train,y_train)
```

Out[81]:

▼ LinearRegression ⓘ ?
(https://scikit-learn.org/1.4/modules/generated/sklearn.linear_model.LinearRegression.html)
LinearRegression()

```
In [83]: rf=RandomForestRegressor()  
rf.fit(X_train,y_train)
```

```
Out[83]: ▼ RandomForestRegressor ⓘ ?  
RandomForestRegressor()  
(https://scikit-learn.org/1.4/modules/generated/sklearn.ensemble.RandomForestRegressor.html)
```

```
In [84]: xg=XGBRegressor()  
xg.fit(X_train,y_train)
```

```
Out[84]: ▼ XGBRegressor ⓘ ?  
colsample_bylevel=None, colsample_bynode=None,  
colsample_bytree=None, device=None, early_stopping_rounds=None,  
enable_categorical=False, eval_metric=None, feature_types=None,  
gamma=None, grow_policy=None, importance_type=None,  
interaction_constraints=None, learning_rate=None, max_bin=None,  
max_cat_threshold=None, max_cat_to_onehot=None,  
max_delta_step=None, max_depth=None, max_leaves=None,  
min_child_weight=None, missing=nan, monotone_constraints=None,  
multi_strategy=None, n_estimators=None, n_jobs=None,  
num_parallel_tree=None, random_state=None, ...)
```

```
In [86]: xgb=GradientBoostingRegressor()  
xgb.fit(X_train,y_train)
```

```
Out[86]: ▼ GradientBoostingRegressor ⓘ ?  
GradientBoostingRegressor()  
(https://scikit-learn.org/1.4/modules/generated/sklearn.ensemble.GradientBoostingRegressor.html)
```

```
In [87]: Y_pred1=lr.predict(X_test)
Y_pred2=rf.predict(X_test)
Y_pred3=xg.predict(X_test)
Y_pred4=xgb.predict(X_test)
```

```
In [89]: from sklearn import metrics
```

```
In [90]: score1=metrics.r2_score(y_test,Y_pred1)
score2=metrics.r2_score(y_test,Y_pred2)
score3=metrics.r2_score(y_test,Y_pred3)
score4=metrics.r2_score(y_test,Y_pred4)
```

```
In [91]: print(score1,score2,score3,score4)
```

```
0.9305603810622383 0.9716419188476204 0.9092836108103018 0.979725911421164
```

```
In [92]: xgb=GradientBoostingRegressor()
xgb_final=xgb.fit(X,y)
```

```
In [93]: import joblib
```

```
In [94]: joblib.dump(xgb_final,"Car_Price_Predictor")
```

```
Out[94]: ['Car_Price_Predictor']
```

```
In [95]: model=joblib.load("Car_Price_Predictor")
```

```
In [96]: xgb_final.save_model("xgb_model.json")
```

```
-----  
AttributeError                                Traceback (most recent call last)  
Cell In[96], line 1  
----> 1 xgb_final.save_model("xgb_model.json")  
  
AttributeError: 'GradientBoostingRegressor' object has no attribute 'save_model'
```

```
In [97]: import pickle
```

```
model_filename = "xgb_model.pkl"  
with open(model_filename, "wb") as file:  
    pickle.dump(xgb_final, file)  
  
print(f"Model saved as {model_filename}")
```

```
Model saved as xgb_model.pkl
```